



Part 1: Measure times of sort algorithms

1. Both algorithms should have $O(n \log n)$ performance, do they?

The STL *sort* algorithm had $O(n \log n)$ performance while the STL *heap_sort* algorithm had slightly better than $O(n \log n)$ performance.

2. How do you know? (Answer these questions in your documentation)

Plotting the results of the sort benchmarks (see the two graphs below), I empirically solved the equation for the trendlines.



Part 2: Add another sort algorithm

1. What algorithm did you pick?


Insertion sort.

2. What is the complexity of the algorithm you picked?

$O(n^2)$ was the expected performance.

3. Does the performance seem to match that performance?


Yes, the algorithm matches $O(n^2)$ performance:




Extra Credit (5 points per algorithm)

I added comb, counting, parallel STL sort, and quick sorts (see included source code) with the following results:



A. Comb sort expected performance is $O(n^2/2^p)$ average, and $O(n \log n)$ best case. My function performs closer to the best case rather than average case:




B. My counting sort $O(n + k)$:



C. STL sort using a parallel (concurrent) algorithm matched expected $O(n \log n)$ performance. However, on arrays with more than 1,000 elements, the algorithm performed on average 20x better than the sequential version. My benchmark computer has only 2 cores. Note the comparison chart below.




D. My quicksort matched expected $O(n \log n)$ performance:



Part 3: Ordered sort values

1. Is the performance significantly different than before?

I benchmarked my quicksort algorithm with random versus ordered data (see included source code). The quicksort with ordered data was *significantly slower*.



2. Why or why not?

When the data is ordered, quicksort selects a pivot value of the largest element in the array. When this happens, one of the partitions contains no elements and the other partition contains $n-1$ elements (all but the pivot). Therefore, quicksort recursively calls the subarray of $n-1$, which results in the series of $O(n-1) + O(n-2) + O(n-3) + \dots + O(1)$, which yields an $O(n^2)$ run time (big-O, worst case).